

Task: Feature Design & Build from a Vague Brief

Scenario

A product manager has handed you the following brief — exactly as written, nothing more:

"We want users to be able to set up alerts so they get notified when something important happens in the world — like breaking news, market movements, natural disasters, that kind of thing. Should work for both email and Slack. Make it flexible enough that we can add more channels later. We need an admin view too."

You may use any technology, language, or framework you choose.

There is no further specification. No wireframes. No schema proposal. No agreement on what "something important" means, where the data comes from, or how events are detected.

Your Task

Using AI agents, take this brief from ambiguity to a working implementation.

Before you start building, define your own plan of attack. Decide what deliverables you think this brief requires, what order to tackle them in, and why. Your plan is itself part of the evaluation — we want to see how you scope and structure work, not just how you execute a predefined checklist.

Then execute against your own plan. As you go, **critically assess what the AI produces**: don't assume outputs are correct just because they look plausible. Surface the checks you run, the things you reject, and the course-corrections you make — that judgment is as important as the work itself.

Suggested Timeframe

Plan for **24 hours**. A partial but well-reasoned solution beats a complete but undocumented one — use the time to think, plan, and show your work, not to maximize output.

A Note on What We're Evaluating

We are not evaluating the solution itself. We are evaluating how you arrived at it.

A working, polished implementation with no visible process tells us very little. What we want to see is how you directed the AI: how you framed the problem, what decisions you made and why, where you course-corrected, and how you validated output before accepting it. The artifacts of that process — plans, decision logs, prompt drafts, notes, intermediate outputs — are the submission. The code is evidence, not the point.

We pay particular attention to **how you evaluated the validity of what the AI produced** — the assumptions you challenged, the hallucinations or shortcuts you

caught, the code you rewrote or rejected, and the sanity checks you applied to generated artifacts.

AI Tooling

Use any AI tool(s) you prefer — there is no required or favored product.

What to Submit

Please share your results with us via a GitHub repository that must contain:

- commits for each major milestone, with meaningful commit messages
- your prompt history — all prompts used during the process
- your plans, working artifacts, and evidence of how you worked — design docs, decision logs, prompt drafts, intermediate outputs, notes, scratch files, screenshots, etc.
- your deliverables — as defined by your own plan